

GPU

Graphics Processing Unit

- A graphics processing unit (GPU), is similar CPU
- Designed specifically for performing the complex mathematical and geometric calculations that are necessary for graphics rendering.

Graphics Processing Unit

- A graphics processing unit (GPU) is a computer chip that performs rapid mathematical calculations, primarily for the purpose of rendering images.
- occasionally called **visual processing unit (VPU)**
- GPU is able to render images more quickly than a CPU because of its parallel processing architecture
- Nvidia introduced the first GPU, the [GeForce 256](#), in 1999
- Others include AMD, Intel and ARM.
- In 2012, Nvidia released a virtualized GPU, which offloads graphics processing from the server CPU in a [virtual desktop infrastructure](#).



Graphics Processing Unit

- GPUs are used in
 - Embedded Systems
 - Mobile phones
 - Personal computers
 - Workstations
 - Game consoles

GPU Vs CPU

- A GPU is tailored for highly parallel operation while a CPU executes programs serially
- For this reason, GPUs have many parallel execution units and higher transistor counts, while CPUs have few execution units and higher clock speeds
- A GPU is for the most part deterministic in its operation
- GPUs have much deeper pipelines (several thousand stages vs 10-20 for CPUs)
- GPUs have significantly faster and more advanced memory interfaces as they need to shift around a lot more data than CPUs



What are GPU's Growth?

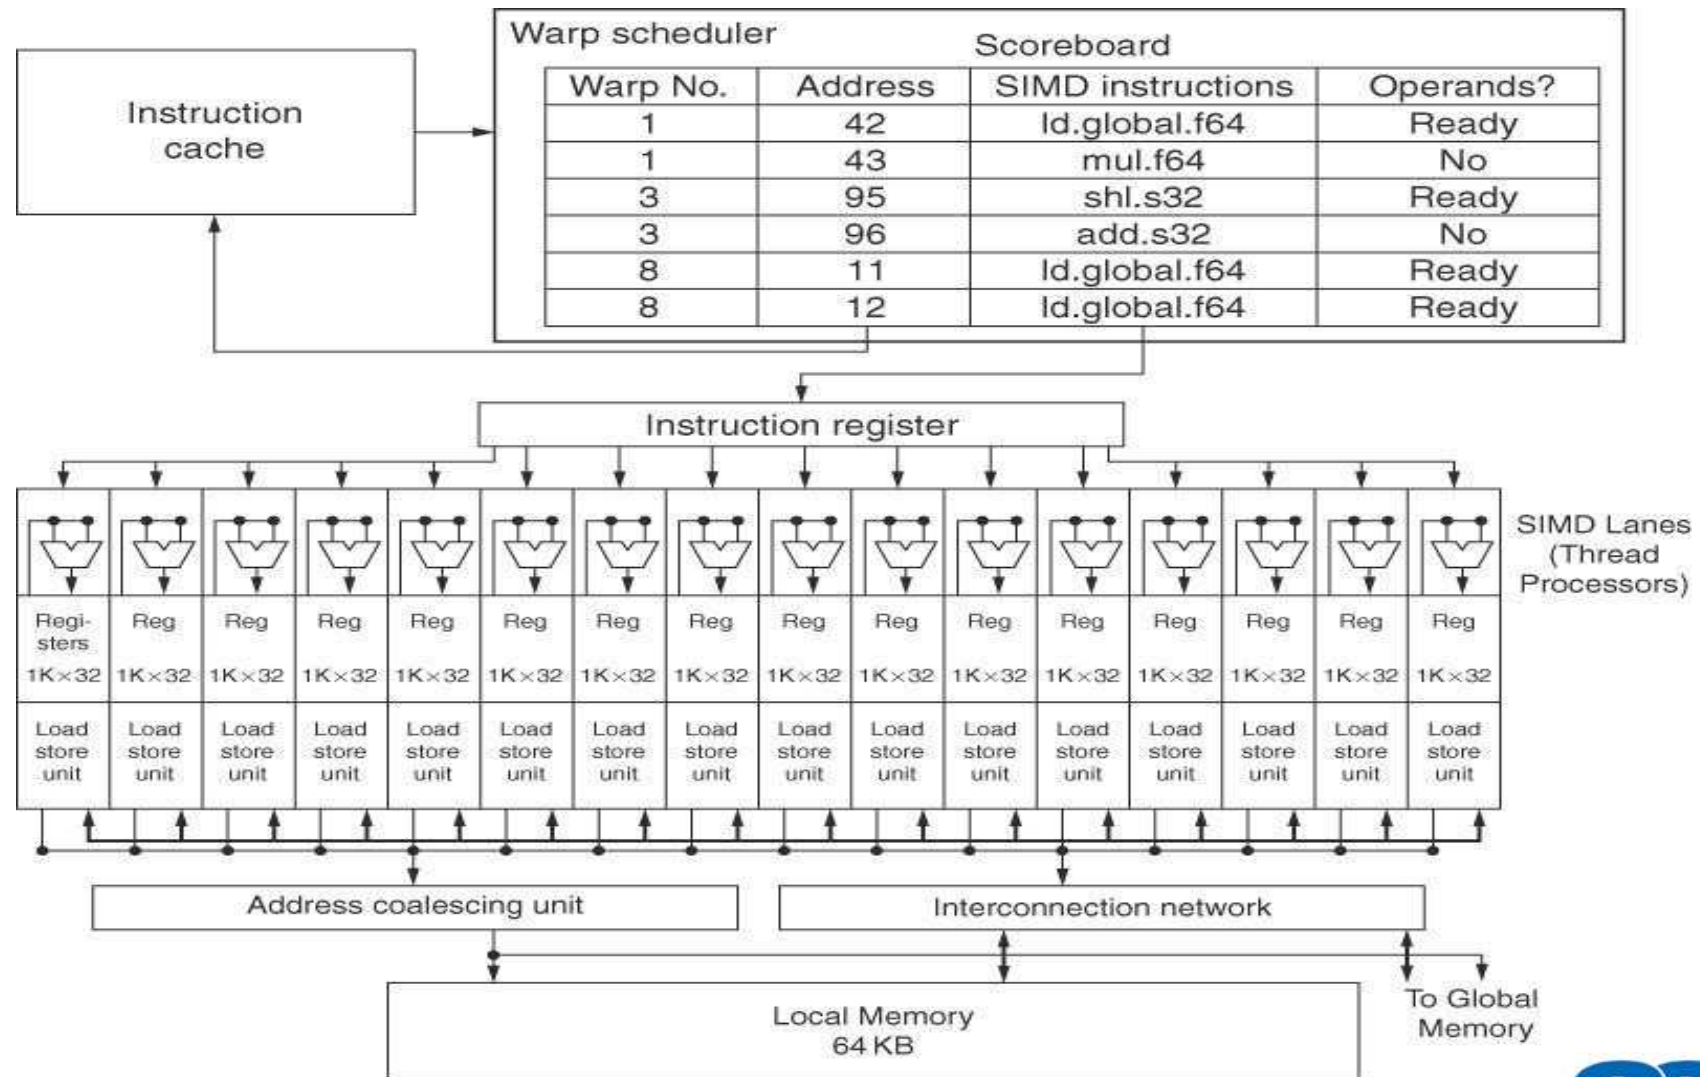
- Entertainment Industry has driven the economy of these chips?
 - Males age 15-35 buy \$10B in video games / year
- Moore's Law ++
- Simplified design (stream processing)
- Single-chip designs



GPU

- Very Efficient For
 - Fast Parallel Floating Point Processing
 - Single Instruction Multiple Data Operations
 - High Computation per Memory Access
- Not Efficient For
 - Double Precision
 - Logical Operations on Integer Data
 - Branching-Intensive Operations
 - Random Access, Memory-Intensive Operations

NVIDIA GPU-MTSIMD



NVIDIA GPU- MTSIMD

- GPU is a multiprocessor composed of MTSIMD processors.
- It is similar to vector processor but with many parallel FU's that are deeply pipelined.
- MTSIMD is a processor that executes code in the form of thread blocks.
- GPU H/W contains a collection of MTSIMD Processors execute a Grid of Thread Blocks.



NVIDIA GPU- MTSIMD

- GPU H/W has two levels of H/W schedulers

1.Thread Block Scheduler:

- Thread block scheduler is similar to control unit in Vector processor
- det the no of thread blocks for a loop and allocates them to diff MTSIMD processors.
- ensures that thread blocks are assigned to the processors whose local memories have the corresponding data.



NVIDIA GPU-FERMI MTSIMD

2. SIMD Thread Scheduler:

- SIMD Thread scheduler has scoreboard logic
- It keeps track of 48 threads of SIMD instructions
- It tells that which thread of SIMD instructions are ready to run
- It sends those instructions to dispatch unit to be run on MTSIMD processor
- within a SIMD Processor, which schedules when threads of SIMD instructions should run



NVIDIA GPU- MTSIMD

- It has many parallel functional units
- SIMD Processors with separate PCs and are programmed using threads.
- Each MTSIMD Processor is assigned 512 elements of the vectors to work on
- SIMD processors have 32,768 registers
- Like vector processor these registers are logically divided across SIMD lanes.



NVIDIA GPU- MTSIMD

- Each SIMD Thread has 64 vector registers of 32 elements with 32 bit each.
- FERMI has 16 physical lanes each contain 2048 registers
- Thread Blocks would contain $512/32 = 16$ SIMD threads.
- Each thread of SIMD instructions in this example compute 32 of the elements of the computation.



NVIDIA GPU- MTSIMD

- GPU applications have so many threads of SIMD instructions that multithreading can
 - hide the latency to DRAM
 - increase utilization of multithreaded SIMD Processors

NVDA GPU ISA

- PTX(Parallel Thread Execution) provides a stable instruction set for GPUs
- H/W instruction set is hidden from the programmer
- PTX instructions describe the operations on a single CUDA thread
- PTX uses virtual registers
- Translation to machine code is performed in software



NVIDIA GPU ISA

- Format of a PTX instruction is **opcode.type d, a, b, c;**
 - where d is the destination operand; a, b, and c are source operands
- Source operands are 32-bit or 64-bit registers or a constant value.
Destinations are registers, except for store instructions.

NVDA GPU ISA

- the operation type is one of the following:

Type	.type Specifier
Untyped bits 8, 16, 32, and 64 bits	. b8, b16, . b32, b64
Unsigned integer 8, 16, 32, and 64 bits	.U8, . U16, U32, u64
Signed integer 8, 16, 32, and 64 bits	.S8, .S16, . S32, S64
Floating Point 16, 32, and 64 bits	.F16, .F32, .F64



Conditional Branching

- Like vector architectures, GPU branch hardware uses internal masks
- Also uses
 - Branch synchronization stack
 - Entries consist of masks for each SIMD lane
 - i.e. which threads commit their results
- Per-thread-lane 1-bit predicate register, specified by programmer