

Concurrency Control Techniques

Overview

- Introduction
- Locking Techniques
- Timestamp Ordering Technique

Introduction

- Need for Concurrency Control
- When operations of different transactions are executed concurrently by interleaving of operations, several problems are associated:
 - Lost Update problem
 - Dirty Read (Temporary Update) problem
 - Incorrect Summary problem

Introduction

- **Purpose of Concurrency Control**
 - To enforce Isolation (through mutual exclusion) among conflicting transactions.
 - To preserve **database consistency** through consistency preserving execution of transactions.
 - To resolve read-write and write-write conflicts.

Introduction

- To overcome various problems associated with concurrency, several techniques were adopted to control it:
 - **Locking Techniques** for Concurrency Control
 - Concurrency Control Based on **Timestamp Ordering**
 - **Multiversion** Concurrency Control Techniques
 - **Validation** Techniques for Concurrency Control

Locking Techniques

- Locking is an operation which secures
 - (a) permission to Read or
 - (b) permission to Write a data item for a transaction.
- Example: `Lock (X)`. Data item X is locked in behalf of the requesting transaction.
- Unlocking is an operation which removes these permissions from the data item.
- Example: `Unlock (X)`. Data item X is made available to all other transactions.
- Lock and Unlock are Atomic operations.

Shared / Exclusive Lock

- Two locks modes (a) shared (read) and (b) exclusive (write).
- **Shared mode:** Shared lock (X).
 - More than one transaction can apply share lock on X for reading its value but *no write lock* can be applied on X by any other transaction.
- **Exclusive mode:** Write lock (X).
 - Only one write lock on X can exist at any time and *no shared lock* can be applied by any other transaction on X.

	Read	Write
Read	Y	N
Write	N	N

Lock Manager

- Lock Manager: Managing locks on data items.
- Lock table: Lock manager uses it to store the identify of transaction locking a data item, the data item, lock mode and pointer to the next data item locked.
- One simple way to implement a lock table is through linked list.

Transaction ID	Data item id	lock mode	Ptr to next data item
T1	X1	Read	Next

Binary Lock

- Database requires that all transactions should be well-formed.
- A transaction is well-formed if:
 - It must lock the data item before it reads or writes to it.
 - It must not lock an already locked data items and it must not try to unlock a free data item.

Binary Lock

lock_item (X):

```
B: if LOCK (X)=0 (* item is unlocked *)  
    then LOCK (X) ← 1 (* lock the item *)  
    else begin  
        wait (until lock (X)=0 and  
            the lock manager wakes up the transaction);  
        go to B  
    end;
```

unlock_item (X):

```
LOCK (X) ← 0; (* unlock the item *)  
if any transactions are waiting  
    then wakeup one of the waiting transactions;
```

Shared / Exclusive Lock

- The following code performs the read operation:

B: if LOCK (X) = “*unlocked*” then

begin LOCK (X) \leftarrow “read-locked”;

no_of_reads (X) \leftarrow 1;

end

else if LOCK (X) \leftarrow “*read-locked*” then

no_of_reads (X) \leftarrow no_of_reads (X) +1

else begin wait (until LOCK (X) = “unlocked” and

the lock manager wakes up the transaction);

go to B

end;

Shared / Exclusive Lock

- The following code performs the [write operation](#):

B: if LOCK (X) = “[unlocked](#)” then

 LOCK (X) ← “write-locked”;

else begin

 wait (until LOCK (X) = “unlocked” and
 the lock manager wakes up the transaction);

 go to B

end;

Shared / Exclusive Lock

The following code performs the [unlock operation](#):

```
if LOCK (X) = "write-locked" then
    begin LOCK (X) ← "unlocked";
           wakes up one of the transactions, if any
    end
else if LOCK (X) ← "read-locked" then
    begin
        no_of_reads (X) ← no_of_reads (X) -1
        if no_of_reads (X) = 0 then
            begin
                LOCK (X) = "unlocked";
                wake up one of the transactions, if any
            end
        end
    end;
end;
```

Lock Conversion

- **Lock upgrade: existing read lock to write lock**

If T_i has a read-lock (X) and T_j has no read-lock (X) ($i \neq j$) then

convert read-lock (X) to write-lock (X)

else

force T_i to wait until T_j unlocks X

- **Lock downgrade: existing write lock to read lock**

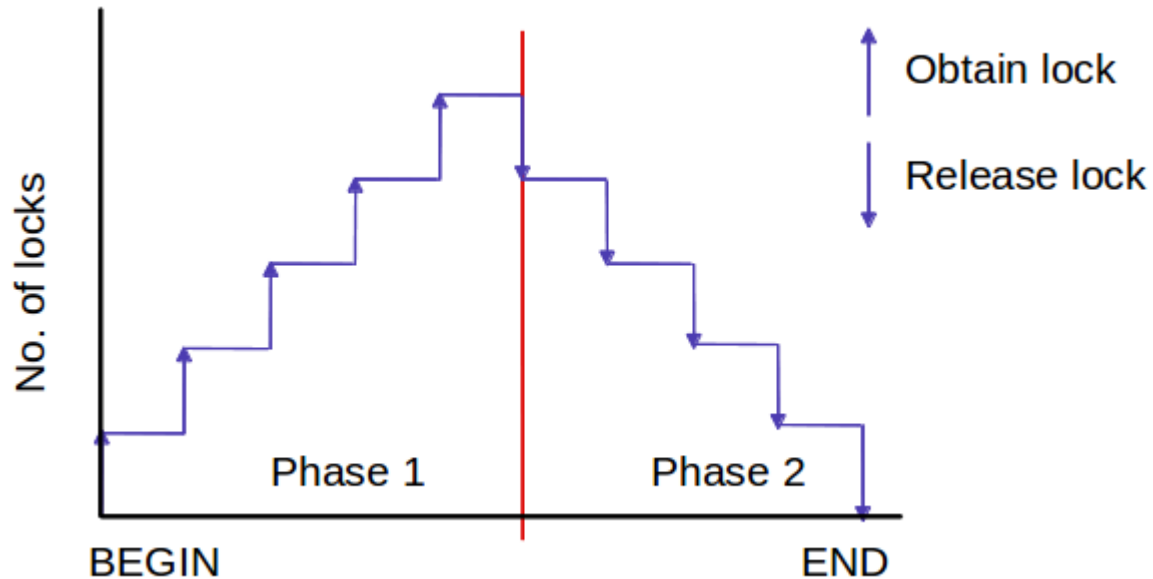
T_i has a write-lock (X) (*no transaction can have any lock on X*)

convert write-lock (X) to read-lock (X)

Two-Phase Locking : Algorithm

- **Two Phases:** (a) Locking (Growing) (b) Unlocking (Shrinking).
- **Locking (Growing) Phase:** A transaction applies locks (read or write) on desired data items one at a time.
 - Locks are *acquired*, **not released**
- **Unlocking (Shrinking) Phase:** A transaction unlocks its locked data items one at a time.
 - Locks are *released*, new locks can **not** be acquired
- **Requirement:** For a transaction these two phases must be mutually exclusive, that is, during locking phase unlocking phase must not start and during unlocking phase locking phase must not begin.

Two-Phase Locking : Algorithm



Two-Phase Locking : Algorithm

(a)

T_1	T_2
<code>read_lock(Y);</code>	<code>read_lock(X);</code>
<code>read_item(Y);</code>	<code>read_item(X);</code>
<code>unlock(Y);</code>	<code>unlock(X);</code>
<code>write_lock(X);</code>	<code>write_lock(Y);</code>
<code>read_item(X);</code>	<code>read_item(Y);</code>
<code>X:=X+Y;</code>	<code>Y:=X+Y;</code>
<code>write_item(X);</code>	<code>write_item(Y);</code>
<code>unlock(X);</code>	<code>unlock(Y);</code>

(b)

Initial values: $X=20, Y=30$

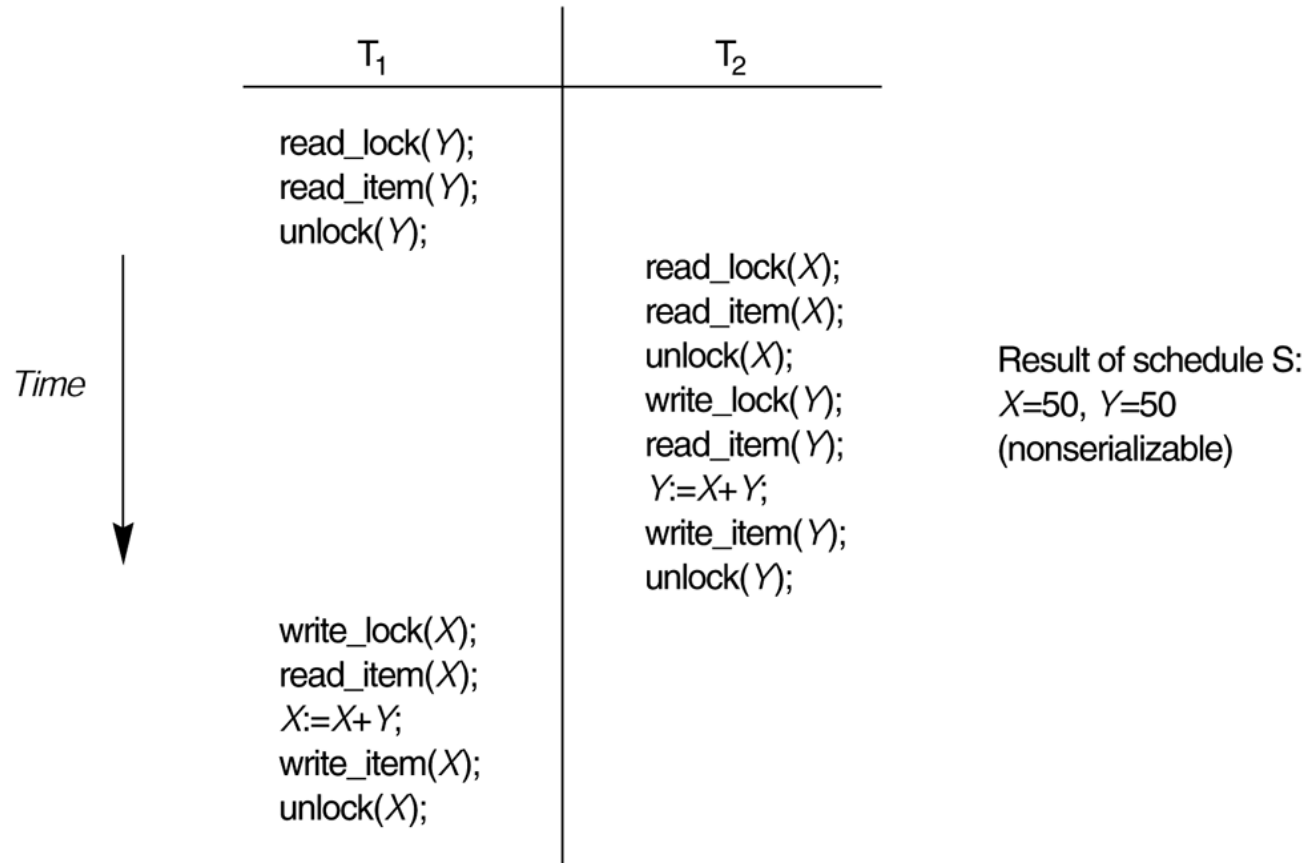
Result of serial schedule T_1 followed by T_2 :
 $X=50, Y=80$

Result of serial schedule T_1 followed by T_2 :
 $X=70, Y=50$

*Transactions T_1 and T_2 that **do not obey** 2PL*

Two-Phase Locking : Algorithm

(c)



A non-serializable schedule that uses locks

Two-Phase Locking : Algorithm

T1 and T2 follow two-phase policy but they are subject to deadlock, which must be dealt with.

T'1

```
read_lock (Y);  
read_item (Y);  
write_lock (X);  
unlock (Y);  
read_item (X);  
X:=X+Y;  
write_item (X);  
unlock (X);
```

T'2

```
read_lock (X);  
read_item (X);  
write_lock (Y);  
unlock (X);  
read_item (Y);  
Y:=X+Y;  
write_item (Y);  
unlock (Y);
```

Two-Phase Locking : Algorithm

- Two-phase policy generates two locking algorithms (a) Basic and (b) Conservative.
- **Basic 2PL:** Transaction locks data items incrementally. This may cause deadlock which is dealt with.
- **Conservative 2PL:** Prevents deadlock by **locking** all desired data items **before transaction begins** execution. (deadlock-free protocol)
- **Strict 2PL:** A more stricter version of Basic algorithm where **unlocking** is performed **after a transaction terminates** (commits or aborts and rolled-back). This is the most commonly used two-phase locking algorithm.

Dealing with Deadlock

Deadlock

T'1

read_lock (Y);
read_item (Y);

write_lock (X);
(waits for X)

T'2

read_lock (X);
read_item (Y);

write_lock (Y);
(waits for Y)

T1 and T2 did follow two-phase policy but they are deadlock

Deadlock (T'1 and T'2)

Dealing with Deadlock

- Deadlock Prevention Protocols
- Deadlock Detection
- Starvation

Dealing with Deadlock

- **Deadlock Prevention Protocols**
- Lock all data items before transaction begins → Conservative 2PL – is a deadlock prevention protocol but further limits concurrency.
- Not generally used because of unrealistic assumptions or overhead.
- What to do with a transaction involved in a deadlock?
 - Should it be blocked and made to wait ?
 - Should it be aborted ?
 - Should the transaction *preempt and abort* another transaction ?
- Transaction timestamp – TS(T) – a unique identifier assigned to each transaction.

Dealing with Deadlock

- **Deadlock Prevention Protocols:**

- If T_i tries to lock an item X ; but X is locked by T_j , then

- **Wait-die**

If $TS(T_i) < TS(T_j)$, then (T_i older than T_j)
 T_i is allowed to wait;

Else

abort T_i and restart it later *with the same timestamp*;

- **Wound-wait**

If $TS(T_i) < TS(T_j)$, then (T_i older than T_j)
abort T_j and restart it later *with the same timestamp*;

Else

T_i is allowed to wait;


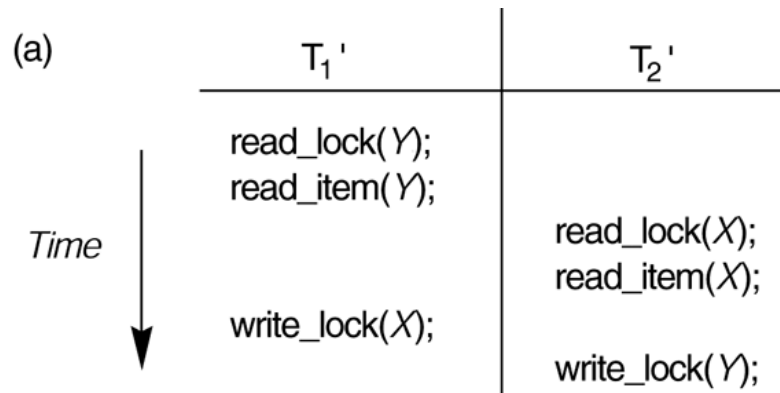
- **Issue:** some transactions to be aborted and restarted needlessly

Dealing with Deadlock

- **Deadlock Detection**

- In this approach, deadlocks are allowed to happen.
- Suitable for transactions that are short and locks only a few items.
- The scheduler maintains a wait-for-graph for detecting cycle. If a cycle exists, then one transaction involved in the cycle is selected (victim) and rolled-back.
- A wait-for-graph is created using the lock table. As soon as a transaction is blocked, it is added to the graph.
- One of the transaction of the cycle is selected and rolled back – select transactions that have not made many changes.

Dealing with Deadlock



Wait-for graph to detect deadlock

Dealing with Deadlock

- **Starvation**
- Starvation occurs when a particular *transaction consistently waits or restarted* and never gets a chance to proceed further.
- In a deadlock resolution it is possible that the same transaction may consistently be selected as victim and rolled-back.
- This limitation is inherent in all priority based scheduling mechanisms.
- Solution: using a *first-come-first-served* queue

Concurrency Control Based on TO


- A timestamp is a unique identifier created by the DBMS to identify a transaction.
- $TS(T)$ refers to timestamp of transaction T .
- Concurrency control techniques based on timestamps **do not use locks**; hence deadlock *cannot* occur.
- Generation of timestamps by counter, system clock
- A larger timestamp value indicates a more recent event or operation.

Concurrency Control Based on TO

- Ordering of transactions are based on their timestamp value – Timestamp Ordering (TO)
- It uses two timestamp TS values:

read_TS(X): the TS of last transaction which reads data item X successfully.

write_TS(X): the TS of last transaction which writes data item X successfully.



Concurrency Control Based on TO

Timestamp based concurrency control algorithm

Basic Timestamp Ordering

1. Transaction T issues a write_item(X) operation:
 - If $\text{read_TS}(X) > \text{TS}(T)$ or if $\text{write_TS}(X) > \text{TS}(T)$, then a younger transaction has already read/written the data item so abort and roll-back T and reject the operation.
 - If the condition in part (a) does not exist, then execute write_item(X) of T and set write_TS(X) to TS(T).
2. Transaction T issues a read_item(X) operation:
 - If $\text{write_TS}(X) > \text{TS}(T)$, then a younger transaction has already written to the data item so abort and roll-back T and reject the operation.
 - If $\text{write_TS}(X) \leq \text{TS}(T)$, then execute read_item(X) of T and set read_TS(X) to the larger of TS(T) and the current read_TS(X).

Concurrency Control Based on TO

- Issues : Basic Timestamp Ordering
- The schedules produced by Basic TO are *conflict serializable*. It leads to cascading rollback problem.
- Basic TO doesn't ensure recoverable schedules; and hence it doesn't ensure cascadeless schedules or *strict schedules*.
- A variation of basic TO called **strict TO** ensures that the schedules are both strict and serializable

Concurrency Control Based on TO

Timestamp based concurrency control algorithm

Strict Timestamp Ordering

1. Transaction T issues a write_item(X) operation:
 - If $TS(T) > read_TS(X)$, then **delay T** until the transaction T' that wrote or read X has terminated (committed or aborted).
2. Transaction T issues a read_item(X) operation:
 - If $TS(T) > write_TS(X)$, then **delay T** until the transaction T' that wrote or read X has terminated (committed or aborted).

Concurrency Control Based on TO

Timestamp based concurrency control algorithm

Thomas's Write Rule

- If $\text{read_TS}(X) > \text{TS}(T)$ then abort and roll-back T and reject the operation.
- If $\text{write_TS}(X) > \text{TS}(T)$, then just ignore the write operation and continue execution. This is because the most recent writes counts in case of two consecutive writes.
- If the conditions given in 1 and 2 above do not occur, then execute $\text{write_item}(X)$ of T and set $\text{write_TS}(X)$ to $\text{TS}(T)$.

References

- Fundamentals of Database Systems, Elmasri and Navathe, 5th Edition



**THANK
YOU!**