

# Recovery Techniques

# Overview

- Recovery Concepts
  - Purpose of Recovery
  - Types of Failure
  - Transaction Log
  - Cascading Rollback
  - Write-Ahead Logging (WAL)
  - Checkpoint
- Deferred Update
- Immediate Update
- Shadow Paging

# Introduction

- **Purpose of Recovery**

- To bring the database into the last consistent state, which existed prior to the failure.
- To preserve transaction properties (Atomicity, Consistency, Isolation and Durability).

# Introduction

- **Types of Failure**
- The database may become unavailable for use due to

Transaction failure: Transactions may fail because of incorrect input, deadlock, incorrect synchronization.

System failure: System may fail because of addressing error, application error, operating system fault, RAM failure, etc.

Media failure: Disk head crash, power disruption, etc.

# Database Recovery

- **System (Transaction) Log**
- For recovery from any type of failure data values prior to modification (BFIM - BeFore Image) and the new value after modification (AFIM – AFter Image) are required.
- These values and other information is stored in a sequential file called Transaction log.
- A sample log is given below. **Back P** and **Next P** point to the previous and next log records of the same transaction.

T ID	Back P	Next P	Operation	Data item	BFIM	AFIM
T1	0	1	Begin			
T1	1	4	Write	X	X = 100	X = 200
T2	0	8	Begin			
T1	2	5	W	Y	Y = 50	Y = 100
T1	4	7	R	M	M = 200	M = 200
T3	0	9	R	N	N = 400	N = 400
T1	5	nil	End			

# Database Recovery

- **Transaction Roll-back (Undo) and Roll-Forward (Redo)**
- To maintain atomicity, a transaction's operations are **redone** or **undone**.

**Undo:** Restore all BFIMs on to disk (Remove all AFIMs).

**Redo:** Restore all AFIMs on to disk.

Database recovery is achieved either by performing only Undos or only Redos or by a combination of the two. These operations are recorded in the log as they happen.

Data	(a)	$T_1$	$T_2$	$T_3$
		read_item( $A$ )	read_item( $B$ )	read_item( $C$ )
		read_item( $D$ )	write_item( $B$ )	write_item( $B$ )
		write_item( $D$ )	read_item( $D$ )	read_item( $A$ )
			write_item( $D$ )	write_item( $A$ )

### Cascading Rollback

(a) The read and write operations of three transactions.

(b) System log at point of crash.

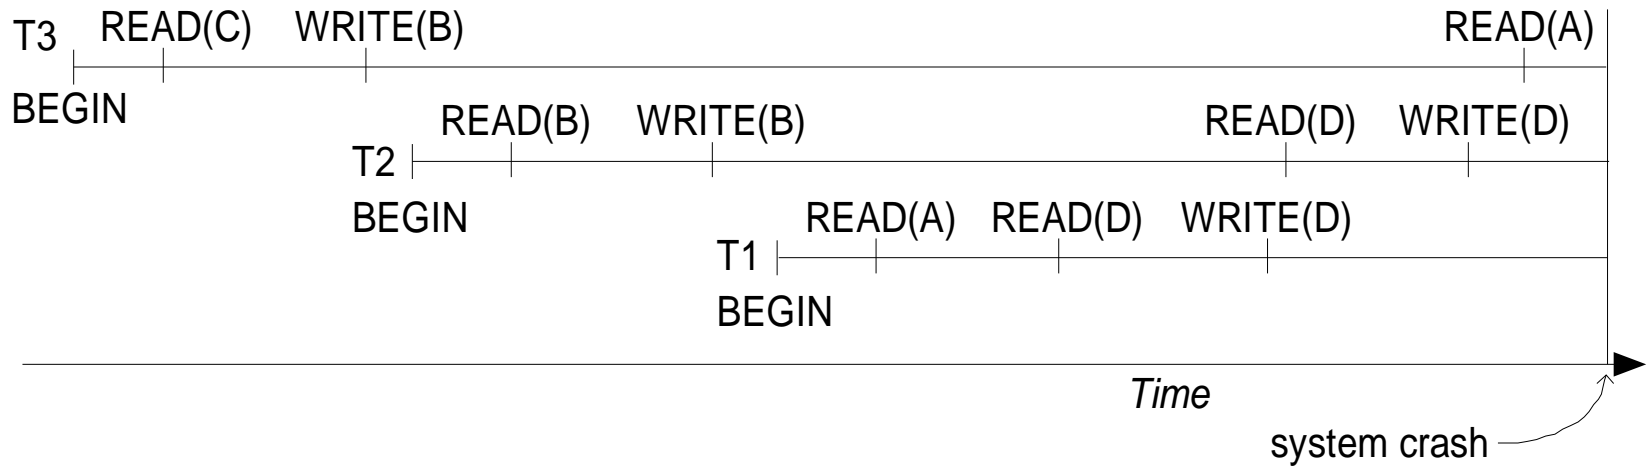
		$A$	$B$	$C$	$D$
		30	15	40	20
(b)	[start-transaction, $T_3$ ]				
	[read_item, $T_3, C$ ]				
*	[write_item, $T_3, B, 15, 12$ ]		12		
	[start-transaction, $T_2$ ]				
	[read_item, $T_2, B$ ]				
**	[write_item, $T_2, B, 12, 18$ ]		18		
	[start-transaction, $T_1$ ]				
	[read_item, $T_1, A$ ]				
	[read_item, $T_1, D$ ]				
	[write_item, $T_1, D, 20, 25$ ]				25
	[read_item, $T_2, D$ ]				
**	[write_item, $T_2, B, 25, 26$ ]				26
	[read_item, $T_3, A$ ]				

← system crash

\*  $T_3$  is rolled back because it did not reach its commit point.

\*\*  $T_2$  is rolled back because it reads the value of item  $B$  written by  $T_3$ .

# Database Recovery



(c) Operations before the crash



# Database Recovery

- **Write-Ahead Logging**

- When **in-place** update (immediate or deferred) is used then log is necessary for recovery and it must be available to recovery manager.
- This is achieved by **Write-Ahead Logging** (WAL) protocol. WAL states:

**For Undo:** Before a data item's AFIM is flushed to the database disk (overwriting the BFIM) its BFIM must be written to the log.

**For Redo:** Before a transaction executes its commit operation, all its AFIMs must be written to the log and the log must be saved on a stable store.

# Database Recovery

- **Checkpointing**
- Time to time (randomly or under some criteria) the database flushes its buffer to database disk to minimize the task of recovery.
- The following steps defines a checkpoint operation:
  1. Suspend execution of transactions temporarily.
  2. Force write modified buffer data to disk.
  3. Write a [checkpoint] record to the log, save the log to disk.
  4. Resume normal transaction execution.

During recovery **redo** or **undo** is required to transactions appearing after [checkpoint] record.

# Database Recovery

- Introduction
- **Deferred Update**
  - Recovery Based on Deferred Update
  - Recovery in Single-user Environment
  - Concurrent Execution in Multi-user Environment
- Immediate Update
- Shadow Paging

# Deferred Update

- Defer or postpone any updates to the database until the transaction completes its execution successfully or reaches its commit point
- Deferred Update (No UNDO/REDO):
  1. A set of transactions records their updates in the log.
  2. At commit point under WAL scheme these updates are saved on database disk.
- No UNDO is required because no AFIM is flushed to the disk before a transaction commits.
- REDO is required in case the system fails after a transaction commits but before all its changes are recorded in the database on disk.

# Deferred Update in Single-User


- (a)
- | $T_1$             | $T_2$             |
|-------------------|-------------------|
| read_item( $A$ )  | read_item( $B$ )  |
| read_item( $D$ )  | write_item( $B$ ) |
| write_item( $D$ ) | read_item( $D$ )  |
|                   | write_item( $D$ ) |
- (b)
- [start-transaction,  $T_1$ ]
  - [write\_item,  $T_1, D, 20$ ]
  - [commit,  $T_1$ ]
  - [start-transaction,  $T_2$ ]
  - [write\_item,  $T_2, B, 10$ ]
  - [write\_item,  $T_2, D, 25$ ] ← system crash

The [write\_item,...] operations of  $T_1$  are redone.  
 $T_2$  log entries are ignored by the recovery process.

# Deferred Update

- Two tables are required for implementing this protocol:
- **Active table:** All active transactions are entered in this table.
- **Commit table:** Committed transactions since the last checkpoint.
- During recovery, all transactions of the **commit** table are redone and all transactions of **active** tables are ignored since none of their AFIMs reached the database.
- WRITE operations of committed transactions are redone *in the order in which they were written to the log*

# Deferred Update Concurrent Users



# Deferred Update - Concurrent Transactions

	$T_1$	$T_2$	$T_3$	$T_4$
(a)	read_item(A) read_item(D) write_item(D)	read_item(B) write_item(B) read_item(D) write_item(D)	read_item(A) write_item(A) read_item(C) write_item(C)	read_item(B) write_item(B) read_item(A) write_item(A)

- (b)
- [start\_transaction,  $T_1$ ]
  - [write\_item,  $T_1, D, 20$ ]
  - [commit,  $T_1$ ]
  - [checkpoint]
  - [start\_transaction,  $T_4$ ]
  - [write\_item,  $T_4, B, 15$ ]
  - [write\_item,  $T_4, A, 20$ ]
  - [commit,  $T_4$ ]
  - [start\_transaction,  $T_2$ ]
  - [write\_item,  $T_2, B, 12$ ]
  - [start\_transaction,  $T_3$ ]
  - [write\_item,  $T_3, A, 30$ ]
  - [write\_item,  $T_2, D, 25$ ] ← system crash

$T_2$  and  $T_3$  are ignored because they did not reach their commit points.

$T_4$  is redone because its commit point is after the last system checkpoint.



# Database Recovery

- Introduction
- Deferred Update
- **Immediate Update**
  - Recovery Based on Immediate Update
  - Recovery in Single-user Environment
  - Concurrent Execution in Multi-user Environment
- Shadow Paging

# Immediate Update

- When a transaction issues an update command, the database can be updated “**immediately**”, without any need to wait for transaction to reach its commit point
- Immediate Update (UNDO/REDO):
  1. Transactions records their updates in the log **before** it is applied to the database using WAL
  2. Transaction is allowed to commit **before** all its changes are written to the database
- UNDO is required because AFIM is flushed to the disk before a transaction commits.
- REDO is required in case the system fails after a transaction commits but before all its changes are recorded in the database on disk.

# Immediate Update

- In a **single-user environment** no concurrency control is required but a log is maintained under WAL.
- Note that at any time there will be one transaction in the system and it will be either in the commit table or in the active table.
- The recovery manager performs:
  1. **Undo** of a transaction if it is in the **active** table.
  2. **Redo** of a transaction if it is in the **commit** table.

# Immediate Update


- In **concurrent execution** environment a concurrency control is required and log is maintained under WAL.
- To minimize the work of the recovery manager checkpointing is used.
- Commit table records transactions to be committed and active table records active transactions.
- During recovery, all transactions of the **commit** table are *redone* and all transactions of **active** tables are *undone*.
- UNDO:
  1. Examine the log entry [**write\_item**, T, X, **old\_value**, **new\_value**], and set the value of X in the database to **old\_value** (BFIM)
  2. UNDO must proceed in the *reverse order* from the order in which the operations were written in the log

# Database Recovery

- Introduction
- Deferred Update
- Immediate Update
- **Shadow Paging**
  - Current directory & Shadow directory
  - No-UNDO/No-REDO Algorithm

# Shadow Paging

- The AFIM does not overwrite its BFIM but recorded at another place on the disk.
- Thus, at any time a data item has AFIM and BFIM (Shadow copy of the data item) at two different places on the disk.



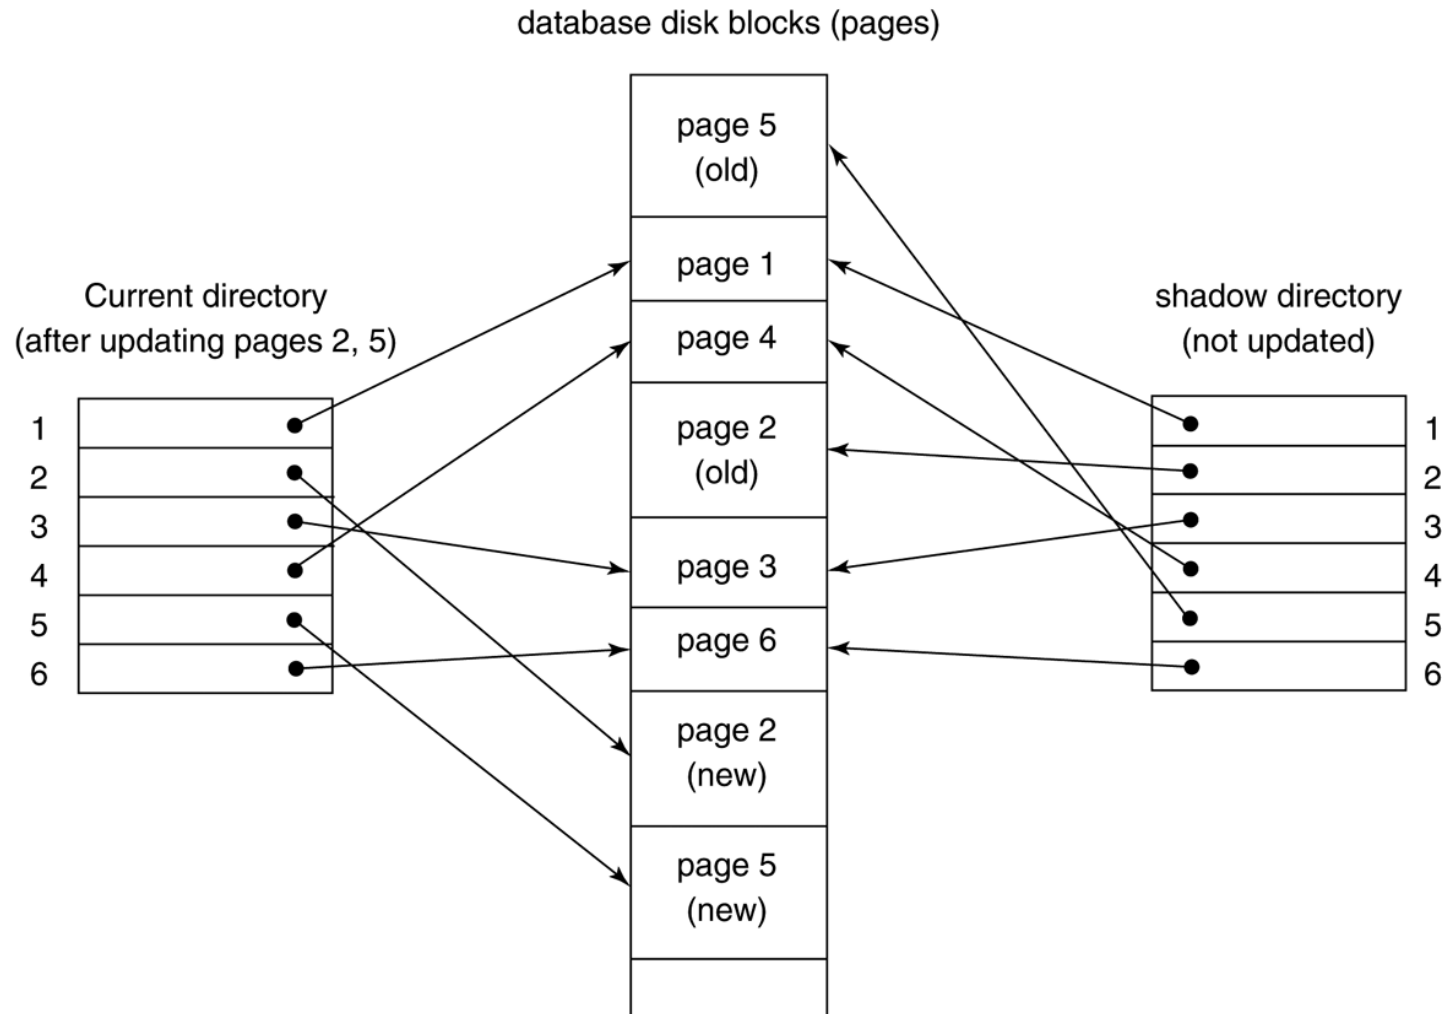
X and Y: Shadow copies of data items

X' and Y': Current copies of data items

# Shadow Paging

- Database is made up of a number of fixed-size disk pages for recovery.
- A directory with  $n$  entries – where  $i^{\text{th}}$  entry points to the  $i^{\text{th}}$  database page on disk.
- When a transaction begins, the **current** directory – points to recent pages on disk – is copied into a **shadow** directory.
- During transaction execution, the shadow directory is *never* modified.
- When write\_item operation, a new copy of database page is created and the current directory entry is modified to point to the new disk block.

# Shadow Paging





# Shadow Paging

- To recover from failure, free the modified database pages and discard the current directory.
- Committing corresponds to discarding the shadow directory.
- Recovery involves neither UNDO nor REDO – **no-undo/no-redo** technique.
- Disadvantages:
  - Need of complex storage management strategies.
  - Overhead of writing shadow directories to disk, if directory is large.
  - Released pages must be added to a list of free pages for future use – garbage collection.

# References

- Fundamentals of Database Systems, Elmasri and Navathe, 3<sup>rd</sup> Edition



**THANK  
YOU!**