

# Schedules

Mirunalini.P

SSNCE

May 20, 2022

# Table of Contents

- 1 Schedules
- 2 Characterizing Schedules Based on Recoverability
- 3 Serializability of Schedules
  - Result Equivalent
  - Conflict Equivalent
- 4 Reference

# Session Objective

- Schedules
- Schedules based on recoverability
- Schedules based on Serializability

At the end of this session, participants will be able to

- Understand the scheduling
- Understand the schedule based on recoverability
- Understand the schedule based on Serializability

**Schedule or History:** A schedule ( or history)  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$  is an **ordering of the operations** of the transactions subject to the constraint that:

- For each transaction  $T_i$  that participates in  $S$ , the operation of  $T_i$  in  $S$  must appear in the same order in which they occur in  $T_i$
- Operations from different transactions can be interleaved in the schedule (such that the operations from transactions  $T_j$  can be interleaved with the operations of  $T_i$  in  $S$ ).

The order of operations in  $S$  is considered to be **total ordering**.

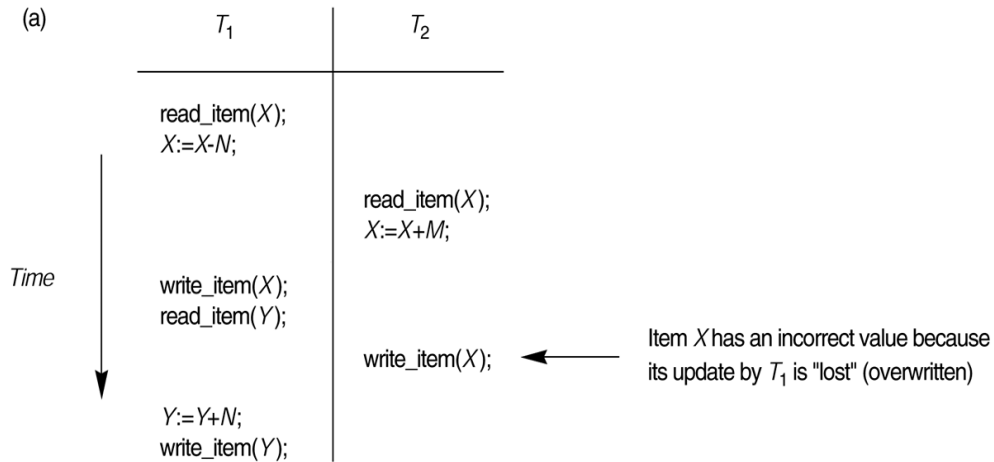
- For any two operations in the schedule, one must occur before the other.

# Schedules

- Schedules can also be displayed in more compact notation
- Order of operations from left to right
- Include only read (r) and write (w) operations, with transaction id (1, 2, ...) and database item name (X, Y, ...)
- Can also include other operations such as b (begin), e (end), c (commit), a (abort)

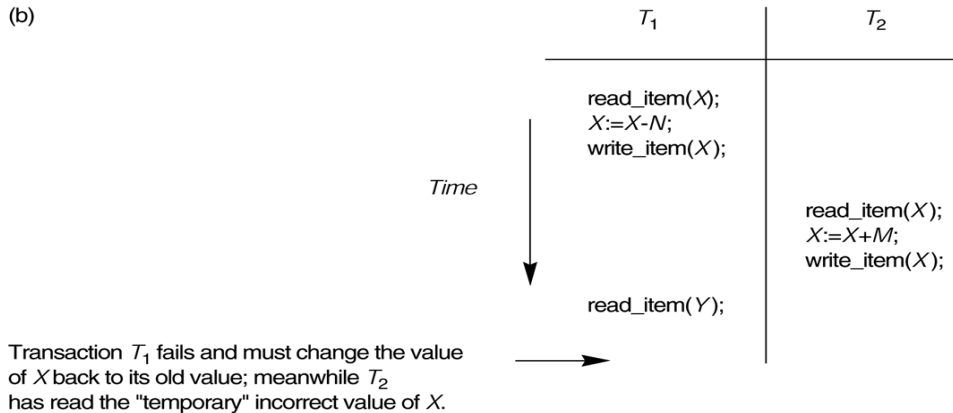
Eg:  $S_a : r1(X); r2(X); w1(X); r1(Y); w2(X); w1(Y);$

# Schedules - Example1



$S_a : r_1(X); r_2(X); w_1(X); r_1(Y); w_2(X); w_1(Y);$

# Schedules - Example2



Assume the transaction  $T_1$  is aborted

$S_b : r_1(X); w_1(X); r_2(X); w_2(X); r_1(Y); a_1$



# Conflicting Operations in a Schedule

Two operations in a schedule are said to be **conflict** if they satisfy all the three conditions:

- if they belong to different transactions
- if they access the same item X
- At least one of the operations is a write\_item(X)

# Conflicting and Non-Conflicting Operations in a Schedule

$S_a : r_1(X); r_2(X); w_1(X); r_1(Y); w_2(X); w_1(Y);$

## Conflicting Operations:

- $r_1(X)$  and  $w_2(X)$
- $r_2(X)$  and  $w_1(X)$
- $w_1(X)$  and  $w_2(X)$

## Non-Conflicting Operations:

- $r_1(X)$  and  $r_2(X)$
- $w_2(X)$  and  $w_1(Y)$
- $r_1(X)$  and  $w_1(X)$

# Conflicting Operations in a Schedule

Two operations conflict if changing their order results in a different outcome

- **Read\_Write Conflict** : Changing  $r_1(X); w_2(X)$  to  $w_2(X); r_1(X)$  means that T1 will read a different value for X
- **Write\_Write Conflict**: Changing  $w_1(Y); w_2(Y)$  to  $w_2(Y); w_1(Y)$  means that the final value for Y in the database can be different
- **Read operations are not conflicting**: Changing  $r_1(Z); r_2(Z)$  to  $r_2(Z); r_1(Z)$  does not change the outcome

# Complete Schedule

A schedule  $S$  of  $n$  transactions  $T_1, T_2, T_3, \dots, T_n$  is said to be **complete schedule** if the following conditions hold:

- The operation in  $S$  are exactly those operations in  $T_1, T_2, T_3, \dots, T_n$  including commit or abort operation as the last operation
- For any pair of operations from the same transaction  $T_i$ , their relative order of appearance in  $S$  is the same as their order of appearance in  $T_i$
- For any two conflicting operations, one of the two must occur before the other in the schedule

Two non-conflicting operations to occur in the schedule without defining which occurs first leads to the **partial order** of the operations in the “ $n$ ” transactions.

# Characterizing Schedules Based on Recoverability

Schedules are characterized as follows:

- **Recoverable schedule:**

- Once a transaction  $T$  is committed, it should never be necessary to roll back ( $T$ )
- This ensures that the durability property of transactions is not violated.
- A schedule  $S$  is **recoverable** if no transaction  $T$  in  $S$  commits until all transactions  $T'$  that have written some item  $X$  that  $T$  reads have committed.

- **Non-Recoverable schedule:**

- A schedule where a committed transaction may have to be rolled back during recovery.
- This violates Durability from ACID properties (a committed transaction cannot be rolled back) and so non-recoverable schedules should not be allowed.

# Characterizing Schedules Based on Recoverability

Consider two schedules:

$Sc : r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); c_2; a_1;$

$Sd : r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); w_1(Y); c_1; c_2;$

**Sc is not recoverable**, because T2 reads item X from T1, and then T2 commits before T1 commits.

If T1 aborts after the C2;

Then the value of X that T2 read is no longer valid and T2 must be aborted after it had already committed — — > **schedule not recoverable**.

For the schedule to be **recoverable** the  $c_2$  operation is postponed until after  $T_1$  commits

**Sd is recoverable**

# Cascading Rollback

An uncommitted transaction has to be rolled back because it read an item from a transaction that failed. This phenomenon is known as **Cascading rollback**

$Se : r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); w_1(Y); a_1; a_2;$

The transaction  $T_2$  has to be rolled back because it read item  $X$  from  $T_1$  which is failed transactions and  $T_1$  then aborted

A schedule is said to be a **Cascadless schedule**, or to avoid cascading rollback, if every transaction in the schedule **reads only items** that were written by committed transactions.

$Se : r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); w_1(Y); a_1; a_2;$

To avoid rollback :

Then  $r_2(X)$  must be postponed untill after  $T_1$  has committed (aborted), thus delaying  $T_2$ .



# Strict Schedule

- A schedule in which a transaction  $T2$  can either **read or write** an item  $X$  until the transaction  $T1$  that last wrote  $X$  has committed
- Simplify the recovery process by simply undoing a `write_item(X)` of an aborted transaction to restore the before image of data item  $X$ .
- **Sf:W1(X,5);W2(X,8);a1**
- if  $T1$  aborts, the recovery procedure recovers the before image which was originally 9, even it has been changed to 8 by transaction  $T2$ .
- $S_f$  is not **strict scheule** since it permits  $T2$  to write item  $X$  even though the last wrote  $X$  had not committed

Schedule C below is **cascadeless and also strict**

**Schedule D is cascadeless, but not strict** (which writes the value of X before T1 commits)

**Schedule E is cascadeless and strict**, to make it strict and cascadeless,  $w_2(X)$  and  $r_2(x)$  must be delayed until after T3 commits

Schedule C:  $r_1(X); w_1(X); r_1(Y); w_1(Y); c_1; r_2(X); w_2(X);$

Schedule D:  $r_1(X); w_1(X); w_3(X); r_1(Y); w_1(Y); c_1; r_2(X); w_2(X);$

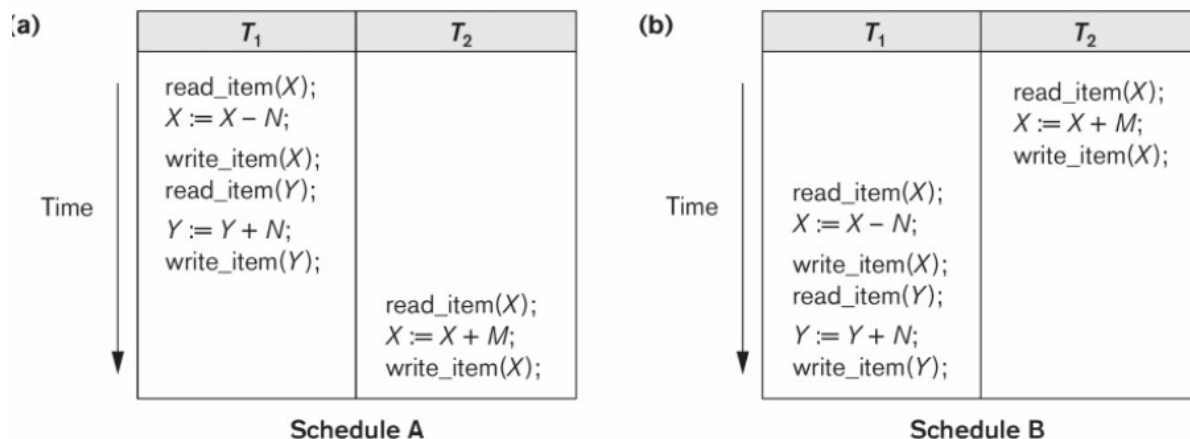
Schedule E:  $r_1(X); w_1(X); r_1(Y); w_1(Y); c_1; w_3(X); r_2(X); w_2(X);$

- The set of all possible schedules can be partitioned into two subsets: **recoverable** and **non-recoverable**
- A **cascadeless** schedule will be the subset of the **recoverable** schedules.
- A **strict schedules** will be the subset of **cascaless** schedules.
- All strict schedules are cascadless and all cascadless schedules are recoverable.

# Serializability of Schedules

- Consider two transactions T1 and T2 which is submitted at the same time.  
If no interleaving is permitted then there are two possible ways:
  - Execute all the operations of T1 and then T2
  - Execute all the operations of T2 and then T1
- **Serializability theory**, attempts to determine which schedules are “correct” and which are not and to develop techniques that allow only correct schedules

# Serial Schedules



**Schedule A and B are called as Serial Schedule**

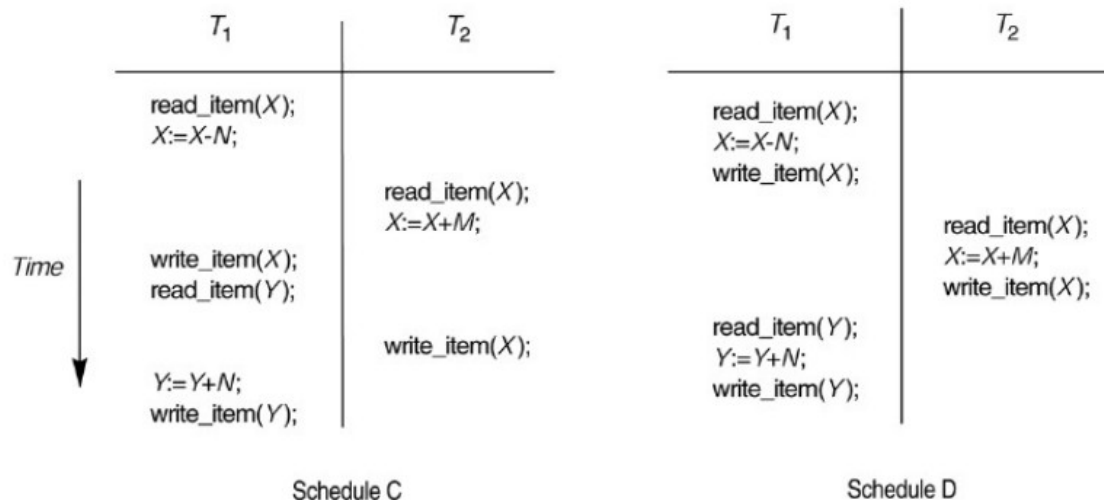
# Serial Schedules

- Schedules A and B are called serial schedules because the operations of each transaction are executed consecutively, without any interleaved operations from the other transaction.
- In a serial schedule:
  - Transactions are performed in serial order
  - Only one transaction is active at any time
  - Commit or abort of the active transactions initiates execution of next transactions
  - Since transactions are independent every serial schedule produce correct result

# Serial Schedules: Disadvantage

- Serial schedules are not feasible for performance reasons: No interleaving of operations.
- Long transactions force other transactions to wait.
- System cannot switch to other transaction when a transaction is waiting for disk I/O or any other event.
- Need to allow concurrency with interleaving without sacrificing correctness

# Non - Serial Schedules

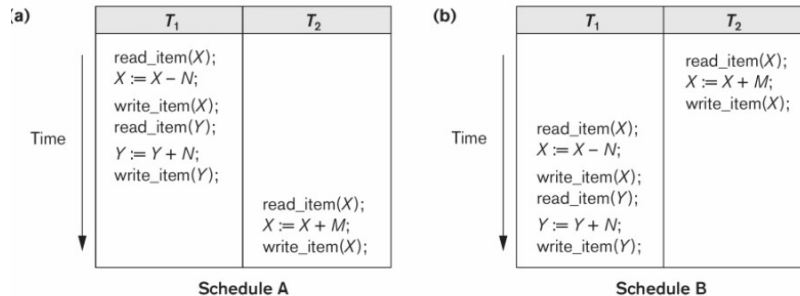
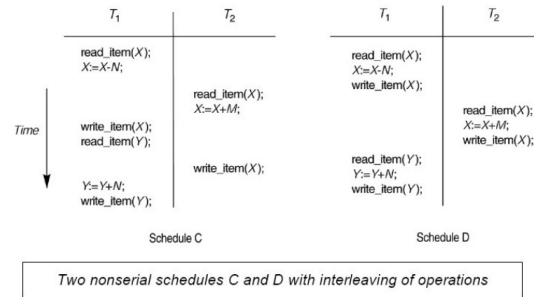


*Two nonserial schedules C and D with interleaving of operations*



# Serial & Non - Serial Schedules

$X=90$  and  $Y=90$ ,  $N=3$  and  $M=2$  serial Schedule A, B expect database values  $X=89$  and  $Y=93$ . C produces the value  $X=92$  and  $Y=93$



# Characterizing Schedule based on Serializability

Schedules C and D are nonserial because of interleaving operations from two transactions.

Some non-serial schedules give correct results.

We have to determine which of the non-serial schedule always give a correct result.

A schedule S of n transactions is **serializable** if it is equivalent to some serial schedule of the same n transactions.

Question is : When are two schedules considered “equivalent” ?

- Result equivalent
- Conflict equivalent - Conflict Serializable
- View equivalent - View Serializable

# Result Equivalent

- Two schedules are called result equivalent if they produce the same final state of the database.
- Disadvantages:
  - Not all the result equivalent schedules produces same final state.
  - Cannot be used to define equivalence of schedules.

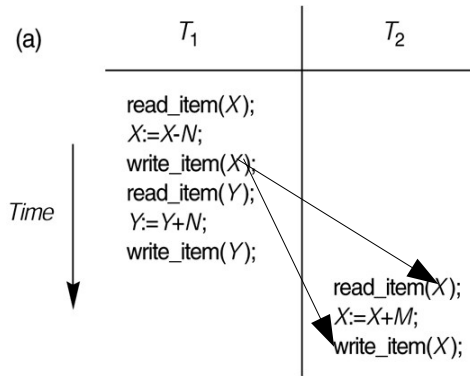
$S_1$	$S_2$
read_item(X);	read_item(X);
$X := X + 10;$	$X := X * 1.1;$
write_item(X);	write_item(X);

*Two schedules that are result equivalent for the initial value of  $X = 100$  but are not result equivalent in general.*

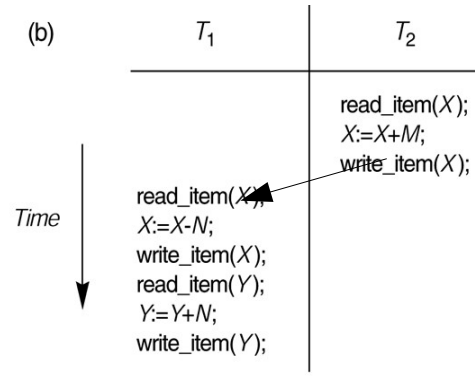
# Conflict Equivalent

- Two schedules are said to be **conflict equivalent** if the order of any two conflicting operations is the same in both schedules.
- According to this, schedule D is equivalent to the serial schedule A. Since A is serial schedule and D is equivalent to A, **D is serializable schedule.**
- We say that a schedule S is **conflict serializable if it is conflict equivalent to a serial schedule**, hence D is conflict serializable.
- Conflict equivalence can be obtained using
  - Ordering of conflict operations
  - Swapping of non-conflict operations
  - Precedence Graph

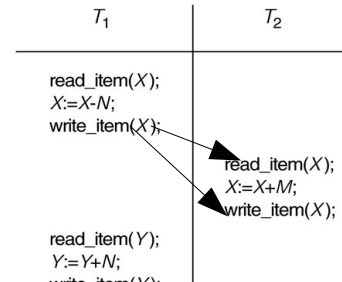
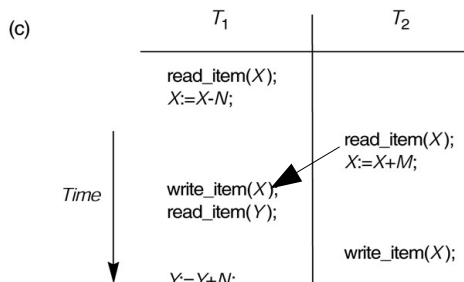
# Conflict Equivalent - Ordering of conflict operations



Schedule A



Schedule B



# Conflict Equivalent - Ordering of conflict operations

$T_1$	$T_2$
read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B)

$S_1$

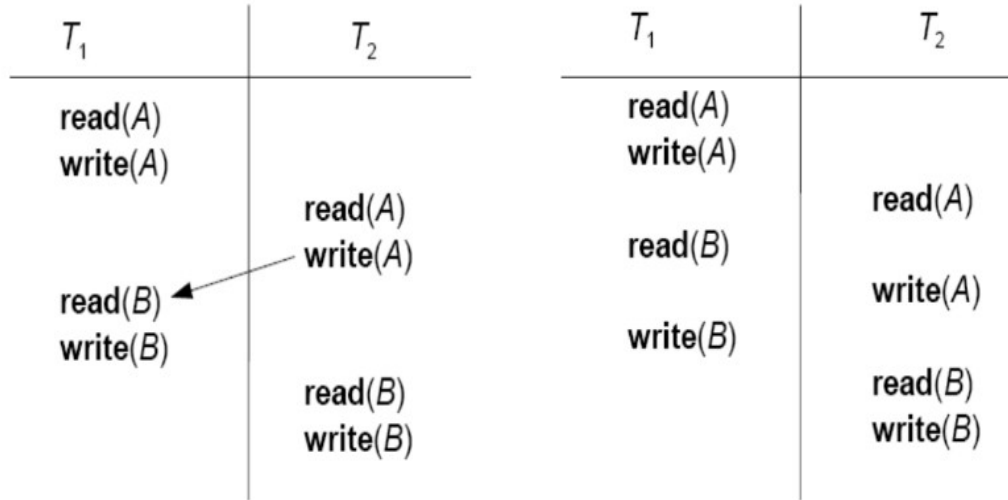
$S_1$

$T_1$	$T_2$
read(A) write(A)  read(B) write(B)	read(A) write(A)  read(B) write(B)

$S_2$

# Conflict Equivalent - Swapping of non-conflict operations

If a schedule S can be transformed into an equivalent schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are conflict serializable.



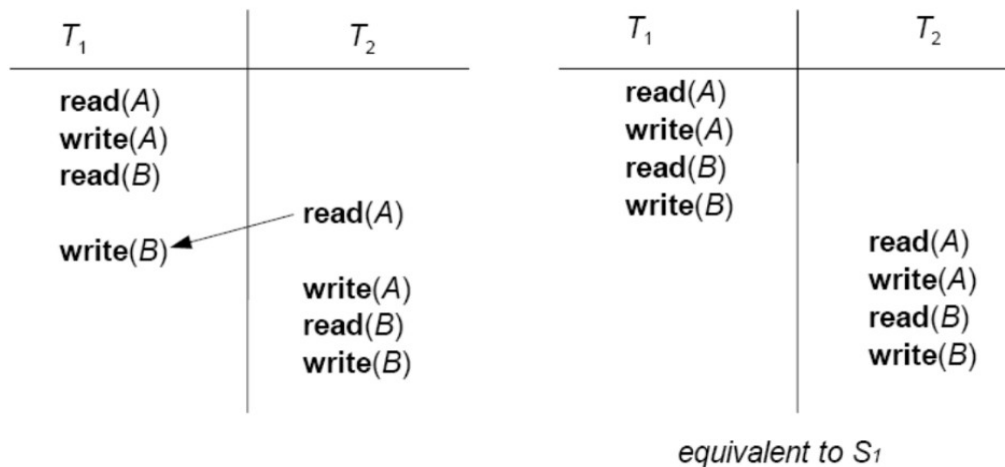
# Conflict Equivalent

$T_1$	$T_2$
read(A)	
write(A)	
read(B)	read(A)
	write(A)
write(B)	read(B)
	write(B)

$T_1$	$T_2$
read(A)	
write(A)	
read(B)	
	read(A)
write(B)	write(A)
	read(B)
	write(B)



# Conflict Equivalent



# Conflict Equivalent - Precedence Graph


- The algorithm looks at only the read\_item and write\_item operations in a schedule to construct a precedence graph (or serialization graph),
- The graph is a directed graph  $G = (N, E)$  that consists of a set of nodes  $N = T_1, T_2, \dots, T_n$  and a set of directed edges  $E = e_1, e_2, \dots, e_m$ .
- There is one node in the graph for each transaction  $T_i$  in the schedule.
- Each edge  $e_i$  in the graph is of the form  $(T_j \rightarrow T_k)$ , where  $T_j$  is the starting node of  $e_i$  and  $T_k$  is the ending node of  $e_i$ .
- Edge from node  $T_j$  to node  $T_k$  is created by the algorithm if a pair of conflicting operations exist in  $T_j$  and  $T_k$
- The conflicting operation in  $T_j$  appears in the schedule before the conflicting operation in  $T_k$ .

# Testing for conflict serializability



- For each transaction  $T_i$  participating in a Schedule S, create a node labeled  $T_i$  in the precedence graph.
- An edge is created from  $T_i$  to  $T_j$  if one of the operations in  $T_i$  appears before a conflicting operation in  $T_j$ .
- The schedule is **serializable** if and only if the **precedence graph** has no cycles.

# Precedence Graph


Serial schedule A



Serial schedule B



Schedule C (not serializable)



Schedule D (serializable, equivalent to A)

# View Equivalent

- A less restrictive definition of equivalence of schedules is called view equivalence.
- A schedule is view serializable if it is view equivalent to a serial schedule.
- Schedules are said to be view equivalent if the following three conditions hold:
  - The same set of transactions participates in S and S' and S and S' includes same operations of those transactions.
  - For any operation,  $R_i(X)$  of  $T_i$  in S, if the value of X read was written by an operation  $W_j(X)$  of  $T_j$  (or if it is the original value of X before the schedule started), the same condition must hold for the value of X read by operation  $R_i(X)$  of  $T_i$  in S'. (Initial read and update read)
  - If the operation  $W_k(Y)$  of  $T_k$  is the last operation to write item Y in S, then  $W_k(Y)$  of  $T_k$  must also be the last operation to write item Y in S'. (Final write)

- **Constrained writes:** The value written by  $w(X)$  in  $T_i$  depends only on the value of  $X$  read by  $r1(X)$  in  $T_i$
- **Blind write:** The value written by  $w(X)$  in  $T_i$  is independent of its old value, so it is not preceded by a read of  $X$  in the transaction  $T$

T1	T2	T3
$r1(x)$		
$w1(x)$		
	$w2(X) - > \textit{Blindwrite}$	
		$w3(X) - > \textit{Blindwrite}$

Table 1: Example -Blind Write



Fundamentals of Database systems 7<sup>th</sup> Edition by Ramez Elmasri.