

CPU Scheduling

Unit-II

Lecture -1

Session Objectives

- ▶ To introduce CPU scheduling, which is the basis for multiprogrammed operating systems
- ▶ To describe various CPU-scheduling algorithms
- ▶ To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system
- ▶ To examine the scheduling algorithms of several operating systems

Session Outcomes

At the end of this session, participants will be able to

- ▶ Discuss various CPU-scheduling algorithms and their evaluation criteria

Agenda

Basic Concepts

Scheduling Criteria

First- Come, First-Served

Shortest-Job-First

Priority

Round Robin

Multilevel Queue

Presentation Outline

Basic Concepts

Scheduling Criteria

First- Come, First-Served

Shortest-Job-First


Priority

Round Robin

Multilevel Queue

Basic Concepts

- ▶ Maximum CPU utilization obtained with multiprogramming
- ▶ **CPU/I/O Burst Cycle**
Process execution consists of a cycle of CPU execution and I/O wait
- ▶ CPU burst followed by I/O burst
- ▶ CPU burst distribution is of main concern



CPU Scheduler

- ▶ Short-term scheduler selects from among the processes in ready queue, and allocates the CPU to one of them
- ▶ Queue may be ordered in various ways
- ▶ CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- ▶ Scheduling under 1 and 4 is nonpreemptive
- ▶ All other scheduling is preemptive
 - ▶ Consider access to shared data
 - ▶ Consider preemption while in kernel mode
 - ▶ Consider interrupts occurring during crucial OS activities

Dispatcher

- ▶ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ▶ switching context
 - ▶ switching to user mode
 - ▶ jumping to the proper location in the user program to restart that program
- ▶ **Dispatch latency** time it takes for the dispatcher to stop one process and start another running

Presentation Outline

Basic Concepts

Scheduling Criteria

First- Come, First-Served

Shortest-Job-First

Priority

Round Robin

Multilevel Queue

Scheduling Criteria

- ▶ **CPU utilization** keep the CPU as busy as possible
- ▶ **Throughput** # of processes that complete their execution per time unit
- ▶ **Turnaround time** amount of time to execute a particular process
- ▶ **Waiting time** amount of time a process has been waiting in the ready queue
- ▶ **Response time** amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

Scheduling Algorithm Optimization Criteria

- ▶ Max CPU utilization
- ▶ Max throughput
- ▶ Min turnaround time
- ▶ Min waiting time
- ▶ Min response time

Presentation Outline

Basic Concepts

Scheduling Criteria

First- Come, First-Served

Shortest-Job-First

Priority

Round Robin

Multilevel Queue

First- Come, First-Served (FCFS) Scheduling

Process	Burst Time
P1	24
P2	3
P3	3

Suppose that the processes arrive in the order: P_1, P_2, P_3 . The Gantt chart for the schedule is



- ▶ Waiting Time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
Average Waiting Time : $(0 + 24 + 27)/3 = 17$

FCFS Scheduling

- Suppose that the processes arrive in the order: P2, P3, P1. The Gantt chart for the schedule is



- Waiting time for P1 = 6; P2 = 0; P3 = 3
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- Convoy effect** - short process behind long process
- Consider one CPU-bound and many I/O-bound processes

Presentation Outline

Basic Concepts

Scheduling Criteria

First- Come, First-Served

Shortest-Job-First

Priority

Round Robin

Multilevel Queue


Shortest-Job-First (SJF) Scheduling

- ▶ Associate with each process the length of its next CPU burst
- ▶ Use these lengths to schedule the process with the shortest time
- ▶ SJF is optimal gives minimum average waiting time for a given set of processes
 - ▶ The difficulty is knowing the length of the next CPU request
 - ▶ Could ask the user

Example of SJF

Process	Burst Time
P1	6
P2	8
P3	7
P4	3

SJF scheduling chart



$$\text{Average waiting time} = (3 + 16 + 9 + 0) / 4 = 7$$

Determining Length of Next CPU Burst

- ▶ Can only estimate the length should be similar to the previous one
- ▶ Then pick process with shortest predicted next CPU burst
- ▶ Can be done by using the length of previous CPU bursts, using exponential averaging
 1. t_n = actual length of n^{th} CPU burst
 2. τ_{n+1} = predicted value for the next CPU burst
 3. $\alpha, 0 \leq \alpha \leq 1$
 4. Define $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$
- ▶ Commonly, α set to $\frac{1}{2}$
- ▶ Preemptive version called **shortest-remaining-time-first**

Examples of Exponential Averaging

- $\alpha = 0$
 - $\tau_{n+1} = \tau_n$
 - Recent history does not count
- $\alpha = 1$
 - $\tau_{n+1} = \alpha t_n$
 - Only the actual last CPU burst counts
- If we expand the formula, we get:
$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots$$
$$+ (1 - \alpha)^j \alpha t_{n-j} + \dots$$
$$+ (1 - \alpha)^{n+1} \tau_0$$
- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- Preemptive* SJF Gantt Chart

- Average waiting time = $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$ msec

Presentation Outline

Basic Concepts

Scheduling Criteria

First- Come, First-Served

Shortest-Job-First

Priority

Round Robin

Multilevel Queue

Priority Scheduling

- ▶ A priority number (integer) is associated with each process
- ▶ The CPU is allocated to the process with the highest priority (smallest integer : highest priority)
 - ▶ Preemptive
 - ▶ Nonpreemptive
- ▶ SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- ▶ Problem: **Starvation** - low priority processes may never execute
- ▶ Solution: **Aging** - as time progresses increase the priority of the process

Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

Priority scheduling Gantt Chart



Average waiting time = 8.2 msec

Presentation Outline

Basic Concepts

Scheduling Criteria

First- Come, First-Served

Shortest-Job-First

Priority

Round Robin

Multilevel Queue


Round Robin (RR)

- ▶ Each process gets a small unit of CPU time (**time quantum q**), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ▶ If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- ▶ Timer interrupts every quantum to schedule next process
- ▶ Performance
 - ▶ q large :FIFO
 - ▶ q small: q must be large with respect to context switch, otherwise overhead is too high

Example of RR with Time Quantum = 4


<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

□ The Gantt chart is:




- Typically, higher average turnaround than SJF, but better **response**
- q should be large compared to context switch time
- q usually 10ms to 100ms, context switch < 10 usec

Time Quantum and Context Switch Time



Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

80% of CPU bursts
should be shorter than q

Presentation Outline

Basic Concepts

Scheduling Criteria

First- Come, First-Served

Shortest-Job-First

Priority

Round Robin

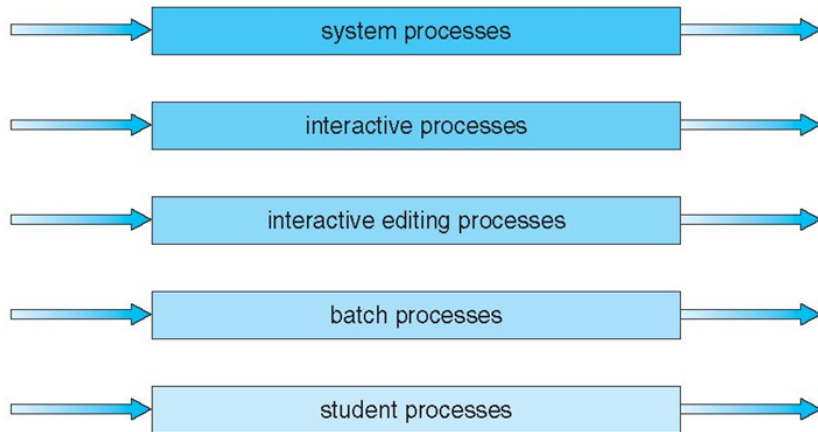
Multilevel Queue

Multilevel Queue

- ▶ Ready queue is partitioned into separate queues, eg:
 - ▶ foreground (interactive)
 - ▶ background (batch)
- ▶ Process permanently in a given queue
- ▶ Each queue has its own scheduling algorithm:
 - ▶ foreground - RR
 - ▶ background - FCFS
- ▶ Scheduling must be done between the queues:
 - ▶ Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - ▶ Time slice - each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - ▶ 20% to background in FCFS

Multilevel Queue Scheduling

highest priority



lowest priority

Multilevel Feedback Queue

- ▶ A process can move between the various queues; aging can be implemented this way
- ▶ Multilevel-feedback-queue scheduler defined by the following parameters:
 - ▶ number of queues
 - ▶ scheduling algorithms for each queue
 - ▶ method used to determine when to upgrade a process
 - ▶ method used to determine when to demote a process
 - ▶ method used to determine which queue a process will enter when that process needs service


Example of Multilevel Feedback Queue

Three queues:

- Q_0 – RR with time quantum 8 milliseconds
- Q_1 – RR time quantum 16 milliseconds
- Q_2 – FCFS


Scheduling

- A new job enters queue Q_0 which is served FCFS
 - ▶ When it gains CPU, job receives 8 milliseconds
 - ▶ If it does not finish in 8 milliseconds, job is moved to queue Q_1
- At Q_1 job is again served FCFS and receives 16 additional milliseconds
 - ▶ If it still does not complete, it is preempted and moved to queue Q_2



Priority-based Scheduling

- ▶ For real-time scheduling, scheduler must support preemptive, priority-based scheduling
- ▶ For hard real-time must also provide ability to meet deadlines
- ▶ Processes have new characteristics: periodic ones require CPU at constant intervals
 - ▶ Has processing time t , deadline d , period p
 - ▶ $0 \leq t \leq d \leq p$
 - ▶ Rate of periodic task is $1/p$




Rate Monotonic Scheduling

- ▶ Rate Monotonic Scheduling assumes that the processing time of a periodic process is the same for each CPU burst
- ▶ static priority policy with preemption
- ▶ A priority is assigned based on the inverse of its period
- ▶ Shorter periods = higher priority;
- ▶ Longer periods = lower priority
- ▶ P1 is assigned a higher priority than P2.


Example

- ▶ Processes, P1 and P2.
- ▶ Periods $p_1 = 50$ and $p_2 = 100$.
- ▶ $t_1 = 20$ for P1 and $t_2 = 35$ for P2.
- ▶ The deadline for each process requires that it complete its CPU burst by the start of its next period.
- ▶ CPU utilization of a process P_i : the ratio of its burst to its period t_i/p_i the CPU utilization of P1 is $20/50 = 0.40$ and that of P2 is $35/100 = 0.35$, **total CPU utilization** : 75



Missed Deadlines with Rate Monotonic Scheduling

- ▶ Assume that process P1 has a period of $p_1 = 50$ and a CPU burst of $t_1 = 25$.
- ▶ For P2, the corresponding values are $p_2 = 80$ and $t_2 = 35$.
- ▶ Rate-monotonic scheduling would assign process P1 a higher priority, as it has the shorter period.
- ▶ The total CPU utilization : $(25/50) + (35/80) = 0.94$



Rate Monotonic Scheduling cannot guarantee that they can be scheduled so that they meet their deadlines

Earliest Deadline First Scheduling (EDF)


- ▶ Priorities are assigned according to deadlines:
- ▶ The earlier the deadline, the higher the priority;
- ▶ The later the deadline, the lower the priority

Earliest Deadline First Scheduling (EDF)

- ▶ EDF scheduling allows process P2 to continue running.
- ▶ P2 now has a higher priority than P1 because its next deadline (at time 80) is earlier than that of P1 (at time 100). Thus, both P1 and P2 meet their first deadlines.
- ▶ Process P1 again begins running at time 60 and completes its second CPU burst at time 85, also meeting its second deadline at time 100.
- ▶ P2 begins running at this point, only to be preempted by P1 at the start of its next period at time 100.
- ▶ P2 is preempted because P1 has an earlier deadline (time 150) than P2 (time 160).
- ▶ At time 125, P1 completes its CPU burst and P2 resumes execution, finishing at time 145 and meeting its deadline.
- ▶ The system is idle until time 150, when P1 is scheduled to run once again.

CPU Scheduling

└ Multilevel Queue



Summary

- ▶ CPU-scheduling algorithms :task of selecting a waiting process from the ready queue and allocating the CPU to it
- ▶ FCFS, SJF, RR, Priority
- ▶ Real time scheduling

Test your understanding

- ▶ Explain the difference between preemptive and nonpreemptive scheduling.
- ▶ What advantage is there in having different time-quantum sizes at different levels of a multilevel queueing system?