


The Linux Operating System

- Krithika Swaminathan

Topics

- Design principles
 - Process Management
 - Scheduling
 - Memory Management
 - File Systems
-




Scheduling

- Job of allocating CPU time to different tasks within an operating system
 - Linux supports preemptive multi-tasking
 - Making decisions – balance fairness and performance
 - Fairness: every process has a chance to get allocated time
 - Performance: best processes or highest priority processes executed
-

Scheduling

- Process Scheduling
 - Real-Time Scheduling
 - Kernel Synchronization
 - Symmetric Multiprocessing
-



Process Scheduling

- Two algorithms:
 - Fair and preemptive
 - Priority-based
 - Completely Fair Scheduler (CFS)
 - Linux Scheduler
 - Terms to remember:
 - Nice value – smaller nice value, higher priority - (-20 to 19)
 - Time slice – length of time the processor is afforded
 - Target latency – interval of time during which every runnable task should run once
 - Minimum granularity – minimum length of time any process should run for
-

Completely Fair Scheduler (CFS)

- Instead of time slices, all processes allotted a proportion of the processor's time
 - Adjusts this allotment by weighting each process's allotment by its nice value
 - Function of the total number of runnable processes
 - N runnable processes --> each afforded $1/N$ of the processor's time
 - Smaller nice value --> receive a higher weight (and vice-versa)
 - (time slice) \propto (process's weight) / (total weight of all runnable processes)
-

Real-Time Scheduling

- Linux implements FCFS and Round Robin
 - Scheduler runs process with the highest priority
 - If equal priority, runs process that has been waiting longest
 - Soft vs Hard real-time scheduling:
 - Hard --> guarantees a minimum latency between when a process becomes runnable and when it really runs
 - Soft --> strict guarantees about relative priorities, but no minimum latency specified
-


Kernel Synchronization

- Request for kernel-mode execution:
 - A running program may request OS service, explicitly or implicitly
 - A device controller may deliver a hardware interrupt
- Problem? - all tasks may try to access same internal DS --> inconsistency
- (critical section problem, shared data)
- Linux kernel provides spinlocks and semaphores for locking in the kernel
- On single-processor machines, spinlocks replaced by enabling and disabling kernel preemption --> `preempt_enable()` and `preempt_disable()`

single processor	multiple processors
Disable kernel preemption.	Acquire spin lock.
Enable kernel preemption.	Release spin lock.

Kernel Synchronization

- Critical sections in interrupt service routines --> interrupt control H/W
- Disabling interrupts --> all I/O suspended --> performance degrades
- Solution - Synchronization architecture – Separating ISR into:
 - Top half:
 - Standard – runs with recursive interrupts disabled
 - Interrupts with same number disabled – others may run
 - Bottom half:
 - Run with all interrupts enabled
 - Invoked automatically when an ISR exits
- kernel can complete any complex processing that has to be done in response to an interrupt without being interrupted itself



Symmetric Multiprocessing

- Linux 2.0 kernel – first stable SMP hardware
 - Separate processes executed in parallel on separate processors
 - Originally, only one processor at a time
 - Version 2.2, single kernel lock (big kernel lock) allowed multiple processes to be active in the kernel concurrently
 - Now, multiple locks – each protects a small subset of kernel's data structures
-

Memory Management

- Management of Physical Memory – pages, blocks of RAM
 - Virtual Memory – memory-mapped into address space of running processes
 - Virtual Memory Regions
 - Lifetime of a Virtual Address Space
 - Swapping and Paging
 - Kernel Virtual Memory
 - Execution and Loading of User Programs
 - Mapping of Programs into Memory
 - Static and Dynamic Linking
-

Management of Physical Memory

- Linux separates physical memory into:
 - ZONE DMA
 - ZONE DMA32
 - ZONE NORMAL
 - ZONE HIGHMEM
- Zones – architecture specific


zone	physical memory
ZONE_DMA	< 16 MB
ZONE_NORMAL	16 .. 896 MB
ZONE_HIGHMEM	> 896 MB

Figure 18.3 Relationship of zones and physical addresses in Intel x86-32.

- Kernel maintains a list of free pages for each zone
-

Management of Physical Memory

- Page allocator
 - responsible for allocating and freeing all physical pages for the zone
 - capable of allocating ranges of physically contiguous pages on request
 - uses a buddy system to keep track of available physical pages
 - Buddy? - adjacent partner of allocatable memory region
 - two allocated partner regions freed up - combined to form larger region - buddy heap
 - converse true - subdivided into partners to satisfy re




Management of Physical Memory

- Memory management sub-systems:
 - k-malloc() variable-length allocator
 - slab allocator, used for allocating memory for kernel data structures
 - page cache, used for caching pages belonging to files
 - Slab
 - used for allocating memory for kernel data structures
 - made up of one or more physically contiguous pages
 - Cache
 - consists of one or more slabs - populated with objects that are instantiations of the kernel DS
 - single cache for each unique kernel DS - ex: file objects, inodes, etc.
-

Management of Physical Memory

- Slab allocation algorithm
- Objects in cache are marked as *free* or *used*
- In Linux, a slab may be in one of three possible states:
 1. Full - All objects in the slab are marked as *used*
 2. Empty - All objects in the slab are marked as *free*
 3. Partial - The slab consists of both *used* and *free* objects



Management of Physical Memory

- Page cache
 - kernel's main cache for files
 - main mechanism through which I/O to block devices performed
 - file systems of all types perform their I/O through the page cache
 - stores entire pages of file contents and is not limited to block devices
-

Virtual Memory

- maintains the address space accessible to each process
- creates pages of virtual memory on demand
- manages loading those pages from disk and swapping them back out to the disk as required
- Logical view
 - address space consists of a set of non-overlapping regions
 - linked into a balanced binary tree to allow fast lookup
- Physical view
 - hardware page tables
 - identify the location of each page of virtual memory, on disk or in physical memory
- vm_area_struct - structure that defines the properties of each region

Virtual Memory

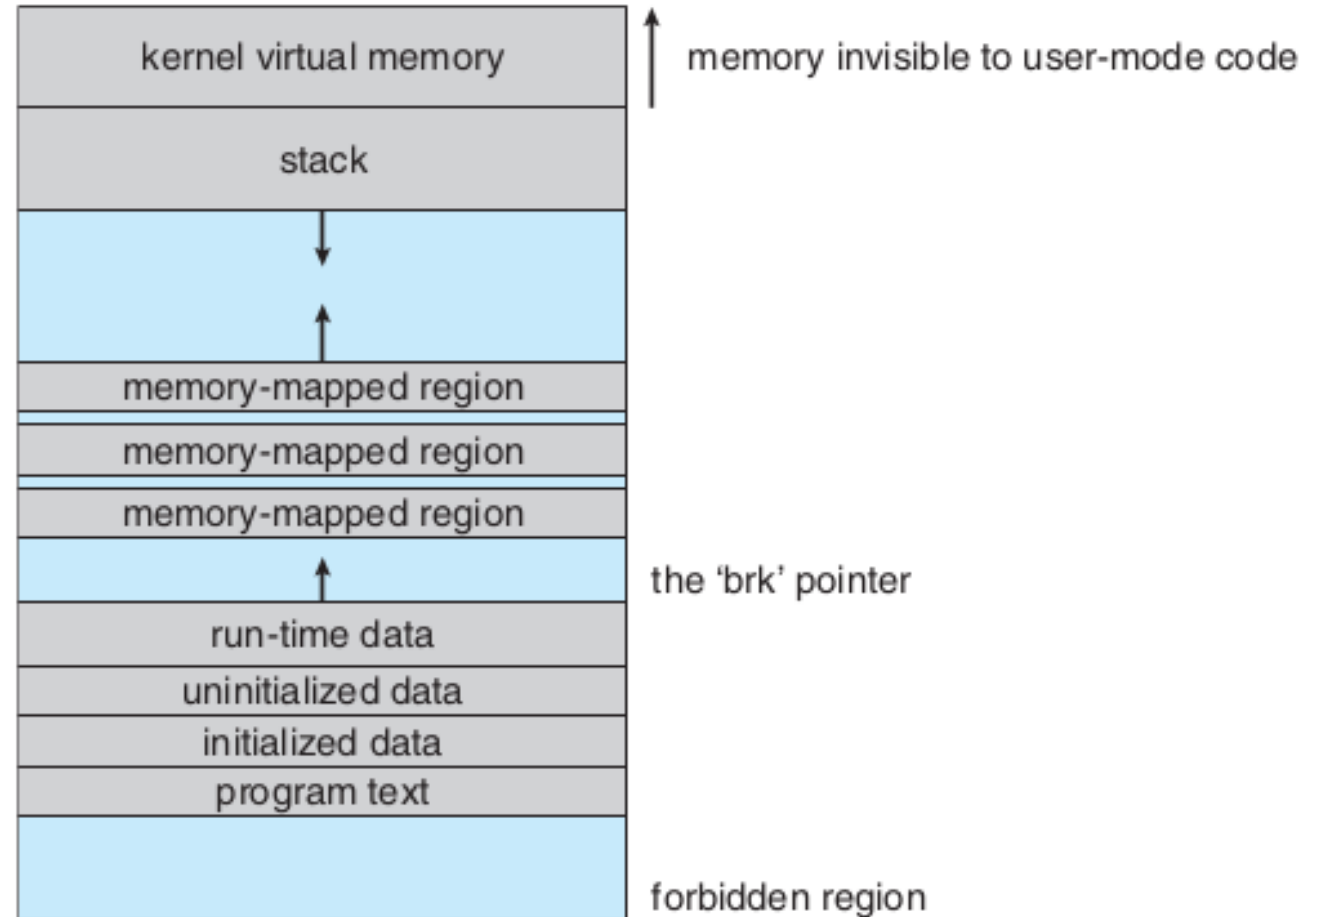
- Virtual Memory region:
 - Backing store; describes where pages come from
 - Demand-zero memory
 - Reaction to writes - private or shared mapping of a region
 - Lifetime of a Virtual Address Space:
 - `exec()`
 - `fork()`
 - Swapping and Paging – paging system:
 - Policy algorithm - decides which pages to write out to disk and when to write them
 - paging mechanism carries out the transfer and pages data back into physical memory
 - Linux's pageout policy – LFU policy (least frequently used)
 - Kernel Virtual Memory – for internal use of Linux
-

Execution and Loading of User Programs

- Older Linux kernels understood a.out format
 - Newer – ELF format
 - Mapping of programs into memory:
 - ELF format binary file consists of a header followed by several page-aligned sections
 - ELF loader works by reading the header and mapping the sections of the file into separate regions of virtual memory
 - Static and Dynamic Linking:
 - Static - necessary library functions embedded directly in the program's executable binary file
 - Dynamic – stub code; contains small, statically linked function for every dynamically linked program
-


Mapping of Programs into Memory

- Kernel – privileged region – inaccessible to normal user-mode programs
- Initialise: stack, program's text and data regions
- Stack at top – grows downward
- Pointer (brk) – points to current extent of data region



File Systems

- The Virtual File System
 - The Linux ext3 File System
 - Journaling
 - The Linux Process File System
-



Virtual File System (VFS)

- The VFS defines four main object types:
 1. An inode object represents an individual file
 2. A file object represents an open file
 3. A superblock object represents an entire file system
 4. A dentry object represents an individual directory entry
 - `int open(. . .)` — Open a file.
 - `ssize_t read(. . .)` — Read from a file.
 - `ssize_t write(. . .)` — Write to a file.
 - `int mmap(. . .)` — Memory-map a file.
-

Virtual File System (VFS)


- inode and file objects are the mechanisms used to access files
 - inode object - a DS containing pointers to the disk blocks that contain the file contents
 - file object - represents a point of access to the data in an open file
 - There is one file object for every instance of an open file, but always only a single inode object.
 - Directory – defines directory or writing data, middle all the tests
 - The superblock object represents a connected set of files that form a self-contained file system.
 - A dentry object represents a directory entry, which may include the name of a directory in the path name of a file
-

Linux ext3

- Each block in a directory file consists of a linked list of entries. In turn, each entry contains the length of the entry, the name of a file, and the inode number of the inode to which that entry refers.
 - The default block size on ext3 varies as a function of the total size of the file system. Supported block sizes are 1, 2, 4, and 8 KB.
 - Allocation policies designed to place logically adjacent blocks of a file into physically adjacent blocks on disk, so that it can submit an I/O request for several disk blocks as a single operation.
 - Block groups, cylinder groups
 - While allocating, try to reduce fragmentation
-

Linux ext3

- Try to keep allocations physically contiguous
- Search for entire free byte; then search for any free bit
- Once free block identified, search extend backward until allocated block encountered
- Reduces CPU cost of disk allocation by allocating multiple blocks simultaneously



Journaling

- Modifications to the file system written sequentially to a journal.
 - When a committed transaction is completed, it is removed from the journal.
 - The journal, a circular buffer, may be in a separate section of the file system, or it may even be on a separate disk spindle.
 - Advantages:
 - Faster performance of operations
 - Updates proceed much faster when they are applied to the in-memory journal rather than directly to the on-disk data structures
 - Why? - performance advantage of sequential I/O over random I/O
-

Linux Process File System

- /proc file system
 - contents are not actually stored anywhere
 - computed on demand according to user file I/O requests
 - Must implement two things:
 - A directory structure
 - File contents within
 - Mapping from inode number to information type - split into two fields:
 - PID
 - type of information being requested about the process
 - maintains a tree data structure of registered global /proc file-system entries
-

SUMMARY

Scheduling:

- Process Scheduling
 - Completely Fair Scheduler
- Real-Time Scheduling
- Kernel Synchronization
- Symmetric Multiprocessing

File Systems:

- The Virtual File System
- The Linux ext3 File System
- Journaling
- The Linux Process File System

Memory Management:

- Management of Physical Memory
- Virtual Memory
 - Virtual Memory Regions
 - Lifetime of a Virtual Address Space
 - Swapping and Paging
 - Kernel Virtual Memory
- Execution and Loading of User Programs
 - Mapping of Programs into Memory
 - Static and Dynamic Linking

THANK YOU!

